

AI Critique

Overall, Claude proved to be a valuable accelerator for API test case generation across FR-06, FR-08, and FR-16, successfully structuring suites covering domain partitions, boundary limits, and general security vectors with well-formed payloads. However, as the test volume scaled, notable inconsistencies and hallucinations emerged: in FR-06, the AI assumed arbitrary database limits ($N = 50$) instead of the actual seed count ($N = 5$), hardcoded brittle error response strings, and generated irrelevant NoSQL injection payloads for an SQLite backend; in FR-08, it relied on non-deterministic token expiry checks, assumed synthetic product prices that conflicted with seed data, and failed to programmatically handle empty-cart state setups; and in FR-16, it generated static ellipsis strings instead of dynamic 5000-character payloads and enforced strict JSON content-type assertions on unrouted GET endpoints. Furthermore, the AI overlooked advanced real-world edge cases like soft-deleted visibility constraints (`is_active=0`), parameter price tampering, mass assignment vulnerabilities, and batch atomicity rollbacks during foreign key violations. To mitigate these failure modes in future workflows, I will refine my prompting strategy by providing explicit schema anchors, system stack constraints (SQLite/Node.js), and strict negative guidelines to prevent requirement hallucination.